0.1 —————————————————————————————————

Scales of measurement

As the previous section indicates, the outcome of a psychological measurement is called a variable. But not all variables are of the same qualitative type and so it's useful to understand what types there are. A very useful concept for distinguishing between different types of variables is what's known as **scales of measurement**.

0.1.1 Nominal scale

A **nominal scale** variable (also referred to as a **categorical** variable) is one in which there is no particular relationship between the different possibilities. For these kinds of variables it doesn't make any sense to say that one of them is "bigger' or "better" than any other one, and it absolutely doesn't make any sense to average them. The classic example for this is "eye colour". Eyes can be blue, green or brown, amongst other possibilities, but none of them is any "bigger" than any other one. As a result, it would feel really weird to talk about an "average eye colour". Similarly, gender is nominal too: male isn't better or worse than female. Neither does it make sense to try to talk about an "average gender". In short, nominal scale variables are those for which the only thing you can say about the different possibilities is that they are different. That's it.

Let's take a slightly closer look at this. Suppose I was doing research on how people commute to and from work. One variable I would have to measure would be what kind of transportation people use to get to work. This "transport type" variable could have quite a few possible values, including: "train", "bus", "car", "bicycle". For now, let's suppose that these four are the only possibilities. Then imagine that I ask 100 people how they got to work today, with this result:

Transportation	Number of people
(1) Train	12
(2) Bus	30
(3) Car	48
(4) Bicycle	10

So, what's the average transportation type? Obviously, the answer here is that there isn't one. It's a silly question to ask. You can say that travel by car is the most popular method, and travel by train is the least popular method, but that's about all. Similarly, notice that the order in which I

list the options isn't very interesting. I could have chosen to display the data like this…

Transportation	Number of people
(3) Car	48
(1) Train	12
(4) Bicycle	10
(2) Bus	30

…and nothing really changes.

0.1.2 Ordinal scale

Ordinal scale variables have a bit more structure than nominal scale variables, but not by a lot. An ordinal scale variable is one in which there is a natural, meaningful way to order the different possibilities, but you can't do anything else. The usual example given of an ordinal variable is "finishing position in a race". You *can* say that the person who finished first was faster than the person who finished second, but you *don't* know how much faster. As a consequence we know that 1st > 2nd, and we know that 2nd > 3rd, but the difference between 1st and 2nd might be much larger than the difference between 2nd and 3rd.

Here's a more psychologically interesting example. Suppose I'm interested in people's attitudes to climate change. I then go and ask some people to pick the statement (from four listed statements) that most closely matches their beliefs:

(1) Temperatures are rising because of human activity

(2) Temperatures are rising but we don't know why

(3) Temperatures are rising but not because of humans

(4) Temperatures are not rising

Notice that these four statements actually do have a natural ordering, in terms of "the extent to which they agree with the current science". Statement 1 is a close match, statement 2 is a reasonable match, statement 3 isn't a very good match, and statement 4 is in strong opposition to current science. So, in terms of the thing I'm interested in (the extent to which people endorse the science), I can order the items as $1 > 2 > 3 > 4$. Since this ordering exists, it would be very weird to list the options like this…

(3) Temperatures are rising but not because of humans

(1) Temperatures are rising because of human activity

(4) Temperatures are not rising

(2) Temperatures are rising but we don't know why

...because it seems to violate the natural "structure" to the question.

So, let's suppose I asked 100 people these questions, and got the following answers:

Response	Number
(1) Temperatures are rising because of human activity	51
(2) Temperatures are rising but we don't know why	20
(3) Temperatures are rising but not because of humans	10
(4) Temperatures are not rising	19

When analysing these data it seems quite reasonable to try to group (1), (2) and (3) together, and say that 81 out of 100 people were willing to *at least partially* endorse the science. And it's *also* quite reasonable to group (2), (3) and (4) together and say that 49 out of 100 people registered *at least some disagreement* with the dominant scientific view. However, it would be entirely bizarre to try to group (1), (2) and (4) together and say that 90 out of 100 people said...what? There's nothing sensible that allows you to group those responses together at all.

That said, notice that while we *can* use the natural ordering of these items to construct sensible groupings, what we *can't* do is average them. For instance, in my simple example here, the "average" response to the question is 1.97. If you can tell me what that means I'd love to know, because it seems like gibberish to me!

0.1.3 Interval scale

In contrast to nominal and ordinal scale variables, **interval scale** and ratio scale variables are variables for which the numerical value is genuinely meaningful. In the case of interval scale variables the *differences* between the numbers are interpretable, but the variable doesn't have a "natural" zero value. A good example of an interval scale variable is measuring temperature in degrees celsius. For instance, if it was 15° yesterday and 18° today, then the 3° difference between the two is genuinely meaningful. Moreover, that 3° difference is *exactly the same* as the 3° difference between 7° and 10°. In short, addition and subtraction are meaningful for interval scale variables.[1]

---

[1]Actually, I've been informed by readers with greater physics knowledge than I that temperature isn't strictly an interval scale, in the sense that the amount of energy required to heat something up by 3° depends on it's current temperature. So in the sense that physicists care about, temperature isn't actually an interval scale. But it still makes a cute example so I'm going to ignore this little inconvenient truth.

However, notice that the $0°$ does not mean "no temperature at all". It actually means "the temperature at which water freezes", which is pretty arbitrary. As a consequence it becomes pointless to try to multiply and divide temperatures. It is wrong to say that $20°$ is *twice as hot* as $10°$, just as it is weird and meaningless to try to claim that $20°$ is negative two times as hot as $-10°$.

Again, lets look at a more psychological example. Suppose I'm interested in looking at how the attitudes of first-year university students have changed over time. Obviously, I'm going to want to record the year in which each student started. This is an interval scale variable. A student who started in 2003 did arrive 5 years before a student who started in 2008. However, it would be completely daft for me to divide 2008 by 2003 and say that the second student started "1.0024 times later" than the first one. That doesn't make any sense at all.

0.1.4 Ratio scale

The fourth and final type of variable to consider is a **ratio scale** variable, in which zero really means zero, and it's okay to multiply and divide. A good psychological example of a ratio scale variable is response time (RT). In a lot of tasks it's very common to record the amount of time somebody takes to solve a problem or answer a question, because it's an indicator of how difficult the task is. Suppose that Alan takes 2.3 seconds to respond to a question, whereas Ben takes 3.1 seconds. As with an interval scale variable, addition and subtraction are both meaningful here. Ben really did take $3.1 - 2.3 = 0.8$ seconds longer than Alan did. However, notice that multiplication and division also make sense here too: Ben took $3.1/2.3 = 1.35$ times as long as Alan did to answer the question. And the reason why you can do this is that for a ratio scale variable such as RT "zero seconds" really does mean "no time at all".

0.1.5 Continuous versus discrete variables

There's a second kind of distinction that you need to be aware of, regarding what types of variables you can run into. This is the distinction between continuous variables and discrete variables. The difference between these is as follows:

- A **continuous variable** is one in which, for any two values that you can think of, it's always logically possible to have another value in between.
- A **discrete variable** is, in effect, a variable that isn't continuous. For a discrete variable it's sometimes the case that there's nothing in the middle.

These definitions probably seem a bit abstract, but they're pretty simple once you see some examples. For instance, response time is continuous. If Alan takes 3.1 seconds and Ben takes 2.3 seconds to respond to a question, then Cameron's response time will lie in between if he took

Table1 The relationship between the scales of measurement and the discrete/continuity distinction. Cells with a tick mark correspond to things that are possible.

	continuous	discrete
nominal		✓
ordinal		✓
interval	✓	✓
ratio	✓	✓

. . . .

3.0 seconds. And of course it would also be possible for David to take 3.031 seconds to respond, meaning that his RT would lie in between Cameron's and Alan's. And while in practice it might be impossible to measure RT that precisely, it's certainly possible in principle. Because we can always find a new value for RT in between any two other ones we regard RT as a continuous measure.

Discrete variables occur when this rule is violated. For example, nominal scale variables are always discrete. There isn't a type of transportation that falls "in between" trains and bicycles, not in the strict mathematical way that 2.3 falls in between 2 and 3. So transportation type is discrete. Similarly, ordinal scale variables are always discrete. Although "2nd place" does fall between "1st place" and "3rd place", there's nothing that can logically fall in between "1st place" and "2nd place". Interval scale and ratio scale variables can go either way. As we saw above, response time (a ratio scale variable) is continuous. Temperature in degrees celsius (an interval scale variable) is also continuous. However, the year you went to school (an interval scale variable) is discrete. There's no year in between 2002 and 2003. The number of questions you get right on a true-or-false test (a ratio scale variable) is also discrete. Since a true-or-false question doesn't allow you to be "partially correct", there's nothing in between 5/10 and 6/10. Table **??** summarises the relationship between the scales of measurement and the discrete/continuity distinction. Cells with a tick mark correspond to things that are possible. I'm trying to hammer this point home, because (a) some textbooks get this wrong, and (b) people very often say things like "discrete variable" when they mean "nominal scale variable". It's very unfortunate.

0.1.6 Some complexities

Okay, I know you're going to be shocked to hear this, but the real world is much messier than this little classification scheme suggests. Very few variables in real life actually fall into these nice neat categories, so you need to be kind of careful not to treat the scales of measurement as if they were

hard and fast rules. It doesn't work like that. They're guidelines, intended to help you think about the situations in which you should treat different variables differently. Nothing more.

So let's take a classic example, maybe *the* classic example, of a psychological measurement tool: the **Likert scale**. The humble Likert scale is the bread and butter tool of all survey design. You yourself have filled out hundreds, maybe thousands, of them and odds are you've even used one yourself. Suppose we have a survey question that looks like this:

> Which of the following best describes your opinion of the statement that "all pirates are freaking awesome"?

and then the options presented to the participant are these:

(1) Strongly disagree
(2) Disagree
(3) Neither agree nor disagree
(4) Agree
(5) Strongly agree

This set of items is an example of a 5-point Likert scale, in which people are asked to choose among one of several (in this case 5) clearly ordered possibilities, generally with a verbal descriptor given in each case. However, it's not necessary that all items are explicitly described. This is a perfectly good example of a 5-point Likert scale too:

(1) Strongly disagree
(2)
(3)
(4)
(5) Strongly agree

Likert scales are very handy, if somewhat limited, tools. The question is what kind of variable are they? They're obviously discrete, since you can't give a response of 2.5. They're obviously not nominal scale, since the items are ordered; and they're not ratio scale either, since there's no natural zero.

But are they ordinal scale or interval scale? One argument says that we can't really prove that the difference between "strongly agree" and "agree" is of the same size as the difference between "agree" and "neither agree nor disagree". In fact, in everyday life it's pretty obvious that they're not the same at all. So this suggests that we ought to treat Likert scales as ordinal variables. On

the other hand, in practice most participants do seem to take the whole "on a scale from 1 to 5" part fairly seriously, and they tend to act as if the differences between the five response options were fairly similar to one another. As a consequence, a lot of researchers treat Likert scale data as interval scale.[*2] It's not interval scale, but in practice it's close enough that we usually think of it as being **quasi-interval scale**.

[*2]Ah, psychology …never an easy answer to anything!

1. 研究デザインについての短い導入

> *実験が終わった後に統計家にアドバイスを求められることは，ほとんどの場合，死後の*
> *検査を求めるようなものである。統計家は実験がなぜ死んだかをいうだけだろう。*
> – Sir Ronald Fisher[1]

　この章では，研究をデザインし，データを集め，そのデータが機能するかどうかチェックする，といったことについての基本的なアイデアについて考えることから始めよう。あなた自身が研究をデザインするのに十分な情報生えられないかもしれないが，ほかの人がやる研究を査定するのに必要な基本的ツールの多くは手に入れられると思う。しかし，この本の狙いはデータの収集よりもデータの分析にあるので，ごくあっさりと全体像を見るにとどめる。この章は，ふたつの意味で「特別」である。まず，ほかの章よりも心理学に特有の状況を扱っている。次に，研究の方法論についての科学的問題にかなりの重点を置いていて，データ分析の統計的問題にはそれほど重点をおいていない。とはいえ，このふたつの問題はお互いに関係しているから，この問題をちょっとばかり詳しく扱うのが，統計のテキストの伝統的なやり方になっている。この章は研究デザインについては **Campbell1963** に，測定の尺度に関する議論については **Stevens1946** に依っている。

1.1

心理学的測定への導入

　データ収集を理解するにはまず，測定について考えなければならない。つまり，ここで我々がやろうとしていることは，人の行動や心についての何かを測るということだ。では「測る」とは何だろう？

---

[1]第 1 回インド統計会議会長講演, 1938. 出典: http://en.wikiquote.org/wiki/Ronald_Fisher

1.1.1 心理学的測定についてのいくつかの考え方

測定そのものは微妙な概念だが，基本的には「なにか」に数字やラベル，あるいはほかのよく定義された記述を割り当てる方法を見つけること，と言える。であるから，次に挙げるようなものも心理学的測定の中に含まれる。

- 私の **年齢** は *33* 歳です。
- 私は**アンチョビが** *好きではない* 。
- 私の **染色体的な性別** は *男性である。*
- 私の **性自認** は *男性である*[*2]

上の短いリストには**太字になっているところ**が "測定されようとしている対象" であり，*イタリックになっているところ* が "測定されたもの" である。じつは，それぞれのケースで生じうる測定全体について考えることで，少し拡張することができる。

- 私の**年齢**は (年単位で数えるから)*0, 1, 2, 3 …,* となる。上限がどうなるかはちょっと曖昧だが，現実的には最大でも *150* ぐらいだと思っていれば良い。それ以上生きた人はいないのだから。
- *アンチョビが好きかどうかは，* 答え方として *好き* あるいは *嫌い，* はたまたどちらでもないとか，ときどきね，と答えるだろう。
- 私の **染色体的な性別**は，ほぼ間違いなく *男性 (XY)* か，*女性 (XX)* だろうが，ごくわずかな例外もある。それはクラインフェルター症候群 *(XXY)* というやつで，ほとんど男性の染色体とおなじである。そのほかにもこうした可能性はあるかもしれない。
- 私の **性自認**は *男性* か *女性* のいずれかだが，私の染色体的な性別とは一致していない。私は自分をどちらでもないとするかもしれないし，はっきりと *トランスジェンダーだ，* と答えるかもしれない。

このように，ある事象 (年齢とか) は取りうる値が明らかのものもあれば，ちょっと特殊なものも

[*2]うむ… ちょっと引っかかるだろ？ このセクションはこの本の最も古いパートだから，ちょっと時代遅れで小っ恥ずかしいものになっている。これが書かれているのは 2010 年で，ここに書いてある事実は *確かに正しい*。これを今の 2018 年にも取って考えると，私はもう 33 歳じゃないし，驚くことじゃない。私の染色体が変わったとは思えないし，私の遺伝子は今も XY 型だとおもう。一方，性自認は… ううん。タイトルページには，今私のことを Daniel じゃなくて Danielle と書いているからわかってもらえるかと思うけど (訳註；Daniel は男性名的表記，Danielle は女性名的表記)，最近は性別に関するアンケートでは男性と答えないようにしていて，むしろ "she/her" で呼ばれる方が嬉しくなっている。話すと長くなるけど！実は，この本ではこのことをどう扱うか少し考えた。この本では作者の声が直接含まれているので，この本での表記を全部Danielle にすると随分違ったものになるんじゃないかと思う。でもそれはかなりの作業量になるから，この本での私の名前は "Dan" で通そうと思う。とにかく Danielle のニックネームとして Dan はバッチリだと思わない？ ちっちゃいことは気にしない。私のことをどう呼んだら良いかわからない読者のために，少しでも気楽になってもらえればと思ってこれを書いてる。あ，アンチョビはまだ好きじゃないけどね (笑)。

ある。しかし，誰かの年齢と言ってもこれよりちょっと微妙なものであることは指摘しておきたい。たとえば，上の例では年齢を一年単位で測定することを想定していた。しかしもしあなたが発達心理学者であったら，これはちょっと雑すぎてて，年齢は*年と月 (月齢)* で表現するかもしれない (もし子供が 2 歳 11 ヶ月であれば，これを "2.11" と書いたりする)。新生児に興味があれば，年齢は生まれてから*何日め*かで表現するだろうし，もしかすると*生後何時間経っ*たかで表現するかもしれない。言い換えると，あなたが測定値に何を許容するかというのが重要なのだ。

　これをもう少し詳しく見てみると，"年齢" という概念は実際まったく正確ではないことに気づかれるかもしれない。一般的に言って，我々が "年齢" というときは，暗に "生まれてからの時間の長さ" を意味している。しかし，それが常に正しいわけではない。たとえば，新生児の目の動きをどのようにコントロールしているかに興味があったとしよう。もしそれぐらい若い子に興味があれば，あなたは "生まれたとき" から始めることが意味のある点だとは思わなくなるかもしれない。アリスが生後 3 週間で，ビアンカが 1 週間後に生まれてきたとして，彼女らに "生まれてから 2 時間後" に出会った時彼女らが "同じ年齢" だったといえるだろうか？ 社会的には，日常生活で年齢について語る時に，生まれた時を起点とするが，それは世界に生まれたって独立した存在として自らを扱ってきた人としての時間の総量，として定義しているからだ。科学的な観点からは，興味があるのはそこだけではない。生物としての人間を考えるなら，概念的には意識を持った生き物として成長してきた過程を考える方が便利だし，その観点から行くとアリスとビアンカは同じ年齢だとは言えない。だからあなたが "年齢" という概念を定義しようとするなら，ふたつのやり方があることになる。意識の時間的長さか，生まれてからの時間的長さか，である。大人を対象にしているときは，その差はほとんど意味をなさないが，新生児の場合はそうではないのだ。

　この問題を超えた，方法論的な問いがある。どの "測定法" が，誰かの年齢を見出すのに用いられるべきだろうか。上で述べたように，そこにはいくつかの可能性がある。

- 単に "あなたは何歳ですか" と尋ねる。この事故報告式は早くて，やすくて，簡単だ。しかしその人が質問を十分理解している場合に限るし，人によっては歳を誤魔化すことがある。
- 権威者 (例えば親御さん) に "あなたのお子さんは何歳ですか？" と尋ねる。この方法は早くて，子供を扱うときには簡単な方法だ。親はいつも周りにいるのだから。でも "意識を持ち始めてから" を考えるときはうまくいかない。というのも，親はその子がいつ意識を持ち始めたかを知る由がないからだ。そのときは，別の権威者 (例えば産科医) に頼る必要があるだろう。
- 公的な記録を調べる。例えば出生証明書とか，死亡証明書である。これは時間はかかるし，手間もかかるが，それなりの利点がある (例えば故人を対象にする場合は。)

1.1.2 操作可能にする: あなたの測定を定義する

上の章で論じだれていたことはすべて，**操作化**という概念に関係している。もう少しこの概念を正確にいうと，操作化とは大きな意味があるが，何かしらぼんやりした概念を，正確な測定に落とし込む手続きのことを指す。この操作化のプロセスは，以下のいくつかの異なる要素を含んでいる。

- あなたが測定しようとしているものを明確にすること。例えば "年齢" も "生まれてからの時間" なのか，"意識が芽生えてからの時間" なのか，どちらがあなたの研究の文脈に沿うものなのか？
- それを測定するのにどの手法にするのか決定すること。事故報告式の測定年齢を使うのか，親に聞くのか，公的な記録を取るのか，どれだろう？ もし事故報告式を使うのなら，どういうフレーズでその質問をするだろう？
- 測定が取りうる値のセットを決定すること。この値は常に数量的なものとは限らないことに注意しよう。測定する年齢は数量的かもしれないが，どんな値が許されているのかをよく考えて見てほしい。年単位の年齢なのか，年月なのか，日，時間，どれだろう？ ほかの測定の種類によっては (例えばジェンダー)，値は数量的なものにならない。しかしすでに述べたように，私たちはどんな値が許されているかを考える必要があるのだ。もし事故報告式でジェンダーを尋ねたら，私たちは相手にどんな選択肢を許すだろうか。単に "男性" か "女性" だけで十分なのか？ "その他" のオプションは必要ないのか？ あるいは，特別な選択肢を考えて，回答者自身の言葉を使うことを許すのか？ そしてもしすべての言語報告を可能な値のセットに含めるのであれば，その回答をどのように解釈するのか？

操作化はトリッキーなビジネスで，"唯一の正解" があるものではない。"年齢" や "性別" のインフォーマルな概念を操作化する方法を選択し，それをフォーマルな測定にすることを選んでも，それはあなたが何を測定したいかに依存するのだ。大抵の場合，あなたの領域で活躍している人たちの科学的コミュニティーが，どうしたらいいかについて十分に確立されたアイデアをすでに持っていることと思う。いいかえると，操作化はケースバイケースで考える必要があると言うことだ。とはいえ，個々人の研究プロジェクトにはそれぞれ特別な問題が他にもいっぱいあるが，一般的な側面もあるものだ。

次に進む前に，用語を整理しておきたい。相互に関連する次の四つのことがある。

- **理論的構成概念**。これはあなたが測定しようとしているもの，例えば "年齢"，"性別"，"ジェンダー" や "意見" といったものだ。理論的な構成概念は直接観測されないし，実際にはちょっと曖昧なものであることが多い。
- **測定**。測定はあなたが客観的にしようとする際につかうツールや手法のこと。調査における質問，行動観察や脳のスキャンなど，すべて測定に数えられる。
- **操作化**。"操作化" という用語は，測定と構成概念の論理的なつながりを指す。あるいは，理

論的構成概念から測定を導き出そうとするプロセスのことである。

- **変数**。最後に，新しい用語を。変数とは，世界に存在する何かに対して，私たちが測定を適用した時の結果である。つまり，変数は実際の "データ" で，最終的に我々のデータセットになるものである。

実際には，科学者でさえこれらの擁護間の区別を曖昧にしようとする傾向があるが，その違いを理解しようとすることはとても有益である。

1.2

測定の尺度水準

前のセクションで示したように，心理学的な測定の結果は変数と呼ばれる。しかしすべての変数が同じ質的なタイプではないので，どういうタイプがあるのかを知っておく必要がある。変数のタイプの違いを区別する大変便利な概念として**測定の尺度水準**というものがある。

1.2.1　名義尺度水準

名義尺度水準の変数は (**カテゴリカルな** (質的な) 変数と言われることもあるが)，異なる取りうる可能性との間に特定の関係がないもののことである。この種の変数は，他の値 "より大きい" とか "より良い" といったいかなる意味も持たないし，その平均を出しても全くなんの意味もない。典型的な例として，"目の色" を考えてみよう。目はブルー，グリーン，ブラウンだったり他の色であったりするが，それのどれかが他のものよりも "大きい" とは言わない。だから，"平均的な目の色" について語るのも憚られる。同様に，ジェンダーもそうだろう。男性が女性よりも良いとか悪いとかいうものではない。同じく "平均的な性別" というのも意味がない。つまり，名義尺度水準の変数は，異なる変数の値は異なるということしか言えない。それだけだ。

もう少し注意深く見てみよう。私が人の通勤・通学手段について研究しているとしよう。私が測定しようとしている変数は，仕事に行くためにどういう種類の移動手段を使うかである。"移動のタイプ" 変数は，それほど多い値を取るものではない。せいぜい "電車"，"バス"，"自家用車"，"自転車" ぐらいだろう。ここでは，この四つの可能性しかないことにしよう。その上で，100 人の人に対して

移動手段	人数
(1) 電車	12
(2) バス	30
(3) 自家用車	48
(4) 自転車	10

どうやって仕事場に行ったかを調べて，次のような結果になったとする。

　さて，移動手段の平均はなんだろう？ 明らかにその答えはひとつではない。愚問としか言いようがない。車で移動することが最もポピュラーな方法で，電車で移動するのがもっともポピュラーでない，それだけである。同時に，私が選択肢を列挙した順番が面白くないことにも注意してほしい。データは次のように表示することだってできるのだ。

移動手段	人数
(3) 自家用車	48
(1) 電車	12
(4) 自転車	10
(2) バス	30

…実質的には何も変わってない。

1.2.2　順序尺度水準

　順序尺度水準にある変数は，名義尺度水準の変数より少し構造化されているが，それほどで大きな変化ではない。順序尺度水準の変数は，異なる状態に自然な意味のある順序を与えたものであるが，それ以上のことはない。よく使われる例としては，"レースの最終順位"である。最初にゴールした人に，2 位の人よりも早かったということはできるが，どれぐらい速かったかということはできない。結果的に，1 位 >2 位ということはわかるし，1 位 >2 位ということもわかるが，1 位と 2 位の差が 2 位と 3 位の差よりも大きいかもしれない。心理学的に面白い例もある。私が気候変動に対する人々の考え方に興味をもっているとしよう。私は，その人の信念にもっとも当てはまる文章がどれかを，(以下にリストした 4 つの文章から) ピックアップして人に尋ねる。

(1) 気温は人間の活動のせいで上昇している
(2) 気温は上昇しているがなぜかはわからない
(3) 気温は上昇しているが人間のせいではない
(4) 気温は上昇していない

　この 4 つの文章は，確かに自然な序列がついている。"現在の科学に対してどれほど同意するか"という意味でだ。文章 1 は明らかに科学的営みにマッチしているし，2 はかなり適合している。3 はそれほどでもないし，4 は現在の科学と逆の立場だ。だから，私の興味のある用語 (人は科学を支持

する度合い) でいうと，項目を 1 > 2 > 3 > 4 の順に並べることができる。この順序づけがある以上，次のように選択肢を並べると奇妙なことになる。

> (3) 気温は上昇しているが人間のせいではない
> (1) 気温は人間の活動のせいで上昇している
> (4) 気温は上昇していない (2) 気温は上昇しているがなぜかはわからない

…これは項目の自然な "構造" に違反しているように見えるからだ。では，100 人にこの質問をして，以下のような答えを得たとしよう。

反応	度数
(1) 気温は人間の活動のせいで上昇している	51
(2) 気温は上昇しているがなぜかはわからない	20
(3) 気温は上昇しているが人間のせいではない	10
(4) 気温は上昇していない	19

このデータを分析するときは，グループ (1),(2),(3) を一緒にするのは合理的だと思われるし，100 人のうち 81 人が少なくとも*一部は*現在の科学を支持していると言えそうだ。また，グループ (2),(3),(4) をまとめることもあり得る話で，100 人のうち 49 人が現在の科学的観点に，少なくともいくらか*同意していない*といえる。しかし，(1),(2) と (4) を一緒にすることはおかしなことで，100 人のうち 90 人がどうこうしてるとは言えない。これらの反応を一緒くたにする合理的な根拠がないからだ。

　つまり，私たちがこれらの項目が，意味のあるグルーピングで自然に序列づけできるということはできるが，その平均を求めるということはできない。例えば，ここでの例でいうと，"平均的な" 反応は 1.97 である。もしその意味がわかるというのなら，ぜひ知りたいのでどうぞ教えてください！

1.2.3　間隔尺度水準

　名義や順序の尺度水準に比べると，間隔尺度水準の変数や比率尺度水準の変数は，かなり意味を持った数値である。感覚尺度水準の変数は，数字の間の*差分*が解釈可能だが，"自然な" ゼロの値をもつものではない。間隔尺度水準の良い例は，セ氏で測られる気温である。例えば，昨日は 15° で今日は 18° だというとき，この 3° の違いに意味がある。もっというと，この 3° の違いは，7° と 10° の間の 3° と全く*同じ*である。つまり，間隔尺度水準では足し算引き算に意味があるのだ[*3]。

　[*3]実は，かなり物理学の知識がある読者に教わったのだが，この温度は厳密な間隔尺度水準ではない。というのも，3° 温度を上げようとするのに必要なエネルギーの量は，その時の気温に依存するらしい。だ k ら，物理学的に注意していうなら，温度は間隔尺度とはいえないのだ。しかし例としていいものなので，このちょっとした不都合な真実は無視しようと思う。

しかし注意してほしいのは，0° が "温度がない" ことを意味するものではない点だ。これが意味するのは "水が凍る温度" であって，これはちょっと恣意的なものだ。だから，温度を掛け算したり割り算したりしようとするのはポイントがずれていることになる。20° は 10° の二倍熱いとは言えないし，20° が −10° のマイナス 2 倍というのも意味がないことだ。

あらためて心理学的な例を見てみよう。大学一年生の態度が時間とともにどう変わっていくかに興味があったとしよう。もちろん，各学生が入学した年を記録するだろう。これは間隔尺度水準の変数だ。2003 年に学生になった人は，2008 に始めた学生よりも 5 年早い。しかし，2008 を 2003 で割って，後者の学生を前者の学生に比べて "1.0024 時間後の人" とは言わないだろう。そんなの全く意味がないじゃないか。

1.2.4 比率尺度水準

四つ目の，そして最後の変数の種類は比率尺度水準の変数といわれ，0 がゼロという意味を持つものであり，掛け算や割り算を許すものである。比率尺度水準にある変数の心理学的な例としては，反応時間 (Response time:RT) がいいだろう。人が問題を解いたり質問に答えたりするのに要する時間を図ることは，様々なタスクでよく見られるものである。というにも，それがタスクがいかに難しいかを示す指標になるからだ。アランが質問に回答するまで 2.3 秒かかり，ベンが 3.1 秒かかったとしよう。感覚尺度水準と同じように，足し算や引き算はどちらもここで意味のある操作である。ベンは実際，3.1 − 2.3 = 0.8 秒長くアランより時間を要した。しかし，今回は掛け算や割り算も意味があるのである。つまり，ベンはアランの 3.1/2.3 = 1.35 倍長く時間がかかったとも言える。比率尺度水準にはなぜこんなことができるかというと，反応時間はまさに "0秒" が "時間が経過していない" ことを表しているからだ。

1.2.5 連続か離散か

あなたが考えておくべき第二の分類方法がある。それはあなたが分析に持ち込もうとしている変数の種類に関するものだ。連続変数と離散変数という区別があるのだ。この区別は以下の通りである。
- 連続変数はどの二つの値を取っても，論理的には必ずその両者の間に値がありうるもの。
- 離散変数とは，連続的でない変数のこと。離散変数は，中間点に対応するものがない。

この定義はおそらくちょっと抽象的なのだが，例を見たらすぐにシンプルなものだとわかってもらえるだろう。たとえば，反応時間は連続変数である。アランが 3.1 秒，そしてベンが 2.3 秒反応に要したとして，キャメロンがその間の 3.0 秒になるかもしれない。デイビッドが 3.031 秒，つまりキャメロンとアランの間に反応時間が来ることもありえる。現実的には，反応時間をそんなに正確に測定することはできないかもしれないが，原則的には確かに可能なのだ。反応時間の 2 つの値の間に，常に新しい値を見つけることができるのであれば，RT は連続変数として考えられる。

Table1.1　尺度水準と離散/連続の関係。チェックマークが入っているセルは可能であることを意味する

	連続	離散
名義		✓
順序		✓
間隔	✓	✓
比率	✓	✓

⋯⋯⋯

　離散変数はこのルールを破った時に現れる。例えば，名義尺度水準は常に離散変数だ。移動手段の種類の例を思い出してみると，電車と自転車の "間" に何かがあるとは言えないし，数学的に 2 と 3 の間に 2.3 があるという意味をもたない。移動手段は離散変数なのだ。同様に，順序尺度水準も常に離散変数である。"2 位" は "1 位" と "3 位" の間に常に入るが，"1 位" と "2 位" の間には論理的に何も生じ得ない。間隔尺度水準と比率尺度水準は事情が違う。上で見たように，反応時間は (比率尺度水準の変数だが) 連続変数である。セ氏で測られる温度も (間隔尺度水準の変数だが) 連続変数である。しかし，学校に通い始めた年 (間隔尺度水準) は離散的である。2002 年と 2003 年の間には X 年が存在しないのだ。正誤判定できる問題の正答数も (比率尺度水準の変数だが)，離散変数である。正誤がわかる質問は，"部分的に正解" とすることができないので，5/10 と 6/10 の間というのが考えられない。表 **??**は尺度水準と連続/離散の区別の関係を要約したものである。チェックマークは可能であることに対応している。強調しておきたいのだが，(a) テキストによっては間違えているものがあって，(b)"離散変数" は "名義尺度水準のこと" と言ったりするのだ。残念なことだ。

1.2.6　複雑なもの

　いいかい，これを聞いたらショックを受けるだろうことはわかるけど，現実世界はこの小さな分類法が提案するよりもずっと厄介なものがある。現実世界においてはこの綺麗なカテゴリーにぴったり当てはまる変数はないので，測定の尺度を堅苦しいルールのように扱わない方がいい。そういうものではないのだ。

　これはガイドラインに過ぎず，違う変数の種類を違うように扱うべき状況について考える助けになるようなものにすぎないのだ。それ以上のものではない。

　古い例を取り上げてみよう。多分ほんとうに古典的な例の，心理学的測定ツールである，リッカートスケールのことを考えてみる。リッカートスケールはあらゆる調査デザインにつかえるものだ。あなた自身，何百，何千回と回答したことがあるだろうし，もしかするとあなた自身も使ったことがあるかもしれない。以下のような調査の質問をしたとする。

"海賊はみんなすごい" という発言に対するあなたの意見を最もよく表しているのは，次のうちどれですか?

そして選択肢が次のように提示される。

(1) 全く同意できない
(2) 同意できない
(3) どちらとも言えない
(4) 同意できる
(5) 強く同意できる

これは 5 件法のリッカート尺度の例で，いくつかの (ここでは 5 つの) 順序づけられた可能性の中から一つを選ぶ。一般的にはそれぞれのケースに言葉で説明が加えられる。しかし，全ての項目が正確に記述されているとは限らない。以下のような表示の仕方も，典型的なリッカートの 5 件法である。

(1) 全く同意できない
(2)
(3)
(4)
(5) 強く同意できる

　リッカートスケールはとても便利なツールだ。ちょっと限定的だとしても。この質問はどの種類の変数になるだろう？　明らかに離散的だ。2.5 に反応することができないのだから。名義尺度水準でないことも明らかだ。順番通りに並んでいるのだから。そして比率尺度水準でないことも明らかだ。自然なゼロがないのだから。

　順序尺度水準か間隔尺度水準なのだろうか？　一つの意見ではあるが，"強く同意できる" と "同意できる" の差分が，"同意できる" と "どちらとも言えない" の差分と同じサイズだということを，証明するのは不可能に思える。実際，日常生活の感覚では，これらは全く同じなはずがない。これに従うと，リッカート尺度は順序尺度水準として扱うべきだということになる。一方で，実践的にはほとんどの参加者が "1 点から 5 点までの尺度" 全体をかなり真面目に扱っているようで，5 つの選択肢の違いが互いに似ているように振る舞う傾向がある。このことから，ほとんどの研究者がリッカート尺度のデータを間隔尺度水準として扱っている[4] これは感覚尺度水準ではないが，実際は十分近似できていると考えて，疑似的間隔尺度水準として扱うのが普通だ。

---

[4]嗚呼，心理学…は何に対しても応えるのが難しいなあ！

測定の信頼性を査定する

　ここまで，理論的な構成概念をどのように操作化し，それによって心理学的測定に変えるかについて，少し考えてきた。そして心理学的測定によって，多くの異なる種類の変数を得る。そこで，次の質問を検討してみよう。測定はうまくいったのか？ これをふたつの関連する用語で表現するなら，*信頼性*と*妥当性*の話になる。単純にいうなら，信頼性とはあなたが測定したものをどれぐらい*正確*に測っているかについてのものであり，妥当性はその測定がどの程度*精度*があったかについてのものである。このセクションでは信頼性について論じよう。妥当性についてはセクション**??**でみていこう。

　信頼性は実にシンプルな概念だ。それは測定の反復可能性，あるいは一貫性を表す言葉である。私の体重を "バスルームの体重計" で測定したものは非常に信頼性が高い。もし私が体重計に乗ったり降りたりを繰り返しても，同じ答えを示し続けるだろう。私の知能について，"母親に尋ねる" のは非常に信頼性が低い。ある非，彼女は私がちょっと賢いと言い，別の日に彼女は私を全くの馬鹿者といった。信頼性の概念は測定が正しいかどうかとは異なるものだ (測定の正しさは妥当性に関係する)。私がバスルームの体重計に乗り降りする時に，ジャガイモの袋を持っていたとしても，信頼性は高いままだ。つまり，同じ答えを出し続けるという意味で。しかし信頼性が高いからといって，私の体重と一致しているかというとそうではなく，間違った値になっている。これを専門的にいうと，*信頼性はあるが妥当ではない*測定ということになる。同様に，私の母が私の知能についていうことは信頼性がないが，彼女がいってることはいくらか正しい。たぶん私はものすごく聡明というわけではないし，彼女が私の知能を推定するときは日によって乱高下するものの，基本的には正しいだろう。これは*信頼性は低いが妥当ではある*測定ということになる。もちろん，もし私の母による推測があまりにも信頼性が低いものであると，彼女が私の知能についていうところの数多くのクレームのうち，どれが実際に正しい表現なのかを見極めることは難しくなるだろう。だから，ある意味で，信頼できない測定というのは実践的な目的において妥当でないものになってしまうのだ。だから，多くの人がいうように，信頼性は妥当性にとっての必要条件 (しかし十分条件ではない) ということになる。

　オーケイ。では信頼性と妥当性の違いがはっきりしたところで，違うやり方で信頼性を測定する方法を考えよう。

- **再検査信頼性**。これは時間が経っても一貫しているかどうかに関するものである。もし後日同じ測定をしたら，同じ答えが得られるだろうか？
- **評定者間信頼性**。これは人が違っても一貫しているかに関するものである。もし誰かが同じ測定をしたら (たとえば，別の人が私の知能について評定したら)，同じ答えが得られるだろうか？
- **平行検査信頼性**。これは理論的に質が等しい測定を使っての一貫性に関するものである。もし

違う体重計を持ってきて私の体重を測定したとして，同じ答えが得られるだろうか？

- **内的一貫性信頼性**。もし測定が同じ機能を持つ異なるいくつものパーツから構成されているとしたら (たとえば，性格検査の質問紙の結果はいくつもの問いを通じて迫っていくものだが)，個々のパーツが同じ答えを出す傾向にあるかどうか。

全ての測定が全ての形式の信頼性を満たすというものではない。たとえば，教育評価は測定の一種として考えることができる。私が教えているある科目，*計算論的認知科学*は，研究課題と試験 (プラス，その他少し) の要素から評価することになっている。試験は研究課題の評価とは幾分異なる側面を評価しようとしているから，評価全体としては内的一貫性が低い。しかし，試験はいくつかの質問から構成されていて，それは同じものを (近似的に) 測定しようとしているから，同じような結果を出す傾向にある。つまり，試験そのものはかなり高い内的一貫性をもっているのだ。これは当然のことである。信頼性を求めるのは，同じものを測定したい場合に限るべきだ!

1.4 ——————————————————————————

変数の "role"：予測変数と結果変数

変数の話から移る前に，最後にもう一つ用語を説明しておこうと思う。普通，我々は研究の結果として多くの変数を手にすることになる。そこで，我々がデータを分析する時に，他の変数に関連づけてある変数を説明しようとすることはよくあることだ。このとき2つの役割，つまり "説明する" と "説明される" という役割の違いを意識することが大事だ。今からこれについて明らかにしていこう。まず，何度も繰り返し参照することになるので，変数を記述するのに数学的な記号を使うことにしよう。"説明される" 変数を Y で，"説明する" 変数を $X_1, X_2,$ のように表記することにする。

X と Y という異なる名前をつけて分析するのは，分析において異なる役割を演じるからだ。これらに対して，古くは独立変数 (Independent Variable,IV)，従属変数 (Dependent Variable,DV) と名付けられていた。The IV is the variable that you use to do the explaining (i.e., X) and the DV is the variable being explained (i.e., Y). 独立変数は説明に使うもの (つまり X) で，従属変数は説明される変数に使うもの (つまり Y) である。この名前の背後にある意味は次のようなものだ。もし X と Y の間に本当に関連があるのなら，Y が X に依存している・従属していると言えるし，研究が "適切に" デザインされていたら，X は他の何物からも独立しているはずである。しかし，個人的にはこのネーミングはマズイと思う。これだと，ミスリーディングであることを忘れそうになるからだ。なぜなら，(a) 独立変数は実際に "何物からも独立である" というわけにはいかないし，(b) 変数間に関係がなかったら，従属変数が独立変数に従属することはないのである。現に，独立変数と従属変数というのがマズいネーミングだと考えるのは私一人ではないので，もっと良さそうな他の名前がいくつもある。この本で使うのは予測変数と結果変数である。このアイデアは，変数 X(予測変数) を使っ

Table1.2　データセットで分析に使われる変数が担う異なる役割について，区別を明確にした用語。この本では古典的な用語は使わず，新しい用語を使うようにしていることに注意してほしい。

変数の役割	古典的な名称	モダンな名称
"説明される"	従属変数 (DV)	結果変数
"説明する'	独立変数 (IV)	予測変数

....

て Y(結果変数) について何らかの推測をしようとするから，というところから来ている[*5]。これについて，表 **??**にまとめておく。

1.5

実験的，あるいは非実験的研究

　あなたが気にするべき，大きな区分基準は "実験的研究" なのか，"非実験的研究" なのか，という違いだ。この区別をする時，本当に話をしているのは，実験者がその研究において人と出来事全体をどの程度コントロールしているか，ということについてである。

1.5.1　実験的研究

　実験的研究の特徴は，研究者が研究のあらゆる側面をコントロールしていること，特に被験者がその研究の途中でどういう経験をするかをコントロールしていることにある。実際的には，研究者は予測変数 (独立変数) を操作したり変化させたりするが，結果変数 (従属変数) は自然に変化するに任せる。このアイデアは，予測変数 (独立変数) を意図的に変化させ，それが結果に何らかの因果的な効果を与えるかどうかを見るためにある。もっというと，予測変数以外に結果変数の原因となるものが存在する可能性がないことを保証するため，他のものは一定に保つか，ある種 "バランスの取れた" やり方で，結果に影響を与えないことを保証するのである。実践上は，他の何ものも実験の結果変数に影響を与えないと考えるのは不可能だし，一定に保つというのも難しい。これについての標準的な解決方法は，**乱打マイゼーション**である。つまり，異なる群にランダムに人を割り付け，群ごとに異なる処置をする (すなわち，違う予測変数の値をもたせる)。ランダマイゼーションについては後ほどもっと詳しく説明するが，ここではランダマイゼーションが群間にみられるあらゆる系統的な変

[*5]しかし腹立たしいことに，別の名前もよく使われる。それをリストアップしようとは思わない－そんなことをしたって何にもならないからだ。ただし，私が "結果変数" としたところが時折 "反応変数" とされることがある，ということだけは指摘しておこう。やれやれ。この種の用語の混乱はとてもよくあることで，まいっちゃうよね。

化の可能性を (無くすことは無理だとしても) 最小化するために行われるものであることを確認できれば十分だろう。

　ごく単純で，まったく非現実的で，非倫理的な例をみてみよう。喫煙が肺がんを引き起こすことを検証したいとする。これを実現する一つの方法は，タバコをお吸う人と吸わない人を見つけ出して，タバコを吸う人の肺がん率が高いかどうかを見ることだろう。これは全く実験的ではない。なぜなら，研究者がタバコを吸う人と吸わない人に対してコントロールしていないからだ。そしてこれは大問題なのである。たとえば，タバコを吸う人は食生活が良くない傾向があるかもしれないし，アスベスト鉱山で働く傾向があるといったことが考えられるからだ。ここでのポイントは，群 (喫煙者と非喫煙者) は多くの点において異なっており，喫煙習慣の違いだけではないということだ。だから喫煙者の方が肺がん率が高いということは，他の何か別の原因があるかもしれず，タバコが原因ではないかもしれないのだ。他の要因のこと (たとえば食習慣とか) は，専門的には "交絡変数" というが，これについてはまた後で話すことにしよう。

　今はとりあえず，適切な実験がどのようなものかを考えたい。ここでの問題点は，喫煙者と非喫煙者は多くの点で異なっているということだった。解決策として，もしあなたに倫理観がなかったとしたら，誰が吸うか吸わないかをコントロールすればよい。具体的には，若い非喫煙者を無作為に 2 つのグループに分けて，その半数を強制的に喫煙者にする。そうすると，半分が喫煙者であるということ以外の点で，量グループが異なるところは非常に低くなる。こうして，もし喫煙者群が非喫煙者群よりも肺がんになる確率が高ければ，(a) 喫煙が癌の原因であり，(b) 我々は殺人者だ，というはっきりした確信を得ることができる。

1.5.2　非実験的研究

　非実験的研究は，広い意味で "研究者が実験で行うようなコントロールをしないあらゆる研究" を指していると言えるだろう。もちろん，科学者はいつもコントロールしたがるものだが，上の例にあったように，コントロールできないとかするべきでない状況というのもある。がんになるかどうかを見るために，人に強制的にタバコを吸わせるのは論理的に大きな問題 (というかほとんど犯罪) だから，実験者がコントロールすべきでない状況の良い例といえるだろう。しかし別の問題もある。倫理的な問題を横に置いたとしても，"喫煙実験" は他に幾つかの問題があるのだ。たとえば，"強制的に" 喫煙させようと言った時，私は非喫煙者のサンプルが喫煙を*始めさせる*という意味で，喫煙者にするという意味だった。これはマッドサイエンティストが好みそうな，確実かつ悪どい実験デザインのように聞こえるが，現実世界での効果を検証する方法としては，それほど健全ではないだろう。たとえば，喫煙が肺がんの原因になるのは，その人の食生活がわるいからで，普段から喫煙をする人は食生活が悪いのかもしれない。しかし我々の実験における "喫煙者" は "普通の" 喫煙者ではないから (つまり，非喫煙者を喫煙者にしたとしても，それは他の普通の，現実の喫煙者が持つ特性を持ってないので)，食生活は良いのかもしれない。そうだとすると，彼らが肺がんになることはないから

我々の研究は失敗だ。これは "普通の" 世界の構造を壊してしまったから (専門的には，これは "人為的な" 結果といわれる) である。

　非実験的な研究をふたつのタイプに分ける区分として注目すべきは，**準実験研究** と **事例研究** である。上で述べた例である，喫煙者と非喫煙者の肺がん発症を検証するために喫煙者と非喫煙者を統制せずに行う研究は，準実験デザインといえる。つまり，これは実験と同じだが，予測変数 (独立変数) を操作していないのである。それでも結果を統計的に分析することはできる。ただし，より注意深く，慎重であらなければならない。

　別のアプローチである事例研究は，一つ，あるいはごく少数の事例をとても詳細に記述することをねらう。一般的にいって，事例研究の結果に統計的な分析を適用することはできないし，わずかな独立した事例から "人は一般に"，という一般的な結論を引き出すことはかなり難しい。とはいえ，事例研究は状況次第で有用なのである。まず，他の手法がない状況。神経心理学では，この問題はよくあることである。特定の脳の領域にダメージを受けているひとをたくさん集めることはできないし，そうするとできることはいくつかの事例について，できるだけ詳細に記述することだけである。しかし，事例研究にも素晴らしい利点はある。多くの人を研究対象にできないからこそ，各事例における特定の要因が果たす役割を理解するために，時間とエフォートをかなり注ぎ込むことができるからだ。これは十分にやる価値のあることである。結果として，事例研究はあなたが実験，準実験デザインで見られるような，より統計学的なアプローチを補うことができる。この本では事例研究についてそれほど多く触れないが，紛れもなく価値のある手法なのである!

1.6

妥当性を検証する

　他の何よりも科学者が欲するものは，研究が "妥当" かどうかである。この**妥当性**という考え方の背後にある考え方は，とてもシンプルなものだ。あなたは自分の研究が信用できるか？ もしできないのなら，それは妥当ではないということだ。しかし，そう言ってしまうのは簡単なのだが，信頼性をチェックするよりも妥当性をチェックすることのほうが難しい。本当のことをいうと，妥当性が実際なんなのかについて正確で明確な同意を得ることはできない。現に，様々な種類の妥当性があって，それらは個々別々の問題を扱っている。また，全ての妥当性が全ての研究に当てはまるというものでもない。ここでは 5 つの異なる妥当性について論じようと思う。

- 内的妥当性
- 外的妥当性
- 構成概念妥当性
- 表面的妥当性
- 生態学的妥当性

First, a quick guide as to what matters here. (1) Internal and external validity are the most important, since they tie directly to the fundamental question of whether your study really works. (2) Construct validity asks whether you're measuring what you think you are. (3) Face validity isn't terribly important except insofar as you care about "appearances". (4) Ecological validity is a special case of face validity that corresponds to a kind of appearance that you might care about a lot.

1.6.1 Internal validity

Internal validity refers to the extent to which you are able draw the correct conclusions about the causal relationships between variables. It's called "internal" because it refers to the relationships between things "inside" the study. Let's illustrate the concept with a simple example. Suppose you're interested in finding out whether a university education makes you write better. To do so, you get a group of first year students, ask them to write a 1000 word essay, and count the number of spelling and grammatical errors they make. Then you find some third-year students, who obviously have had more of a university education than the first-years, and repeat the exercise. And let's suppose it turns out that the third-year students produce fewer errors. And so you conclude that a university education improves writing skills. Right? Except that the big problem with this experiment is that the third-year students are older and they've had more experience with writing things. So it's hard to know for sure what the causal relationship is. Do older people write better? Or people who have had more writing experience? Or people who have had more education? Which of the above is the true *cause* of the superior performance of the third-years? Age? Experience? Education? You can't tell. This is an example of a failure of internal validity, because your study doesn't properly tease apart the *causal* relationships between the different variables.

1.6.2 External validity

External validity relates to the **generalisability** or **applicability** of your findings. That is, to what extent do you expect to see the same pattern of results in "real life" as you saw in your study. To put it a bit more precisely, any study that you do in psychology will involve a fairly specific set of questions or tasks, will occur in a specific environment, and will involve participants that are drawn from a particular subgroup (disappointingly often it is college students!). So, if it turns out that the results don't actually generalise or apply to people and situations beyond the ones that you studied, then what you've got is a lack of external validity.

The classic example of this issue is the fact that a very large proportion of studies in psychology will use undergraduate psychology students as the participants. Obviously, however, the researchers

don't care *only* about psychology students. They care about people in general. Given that, a study that uses only psychology students as participants always carries a risk of lacking external validity. That is, if there's something "special" about psychology students that makes them different to the general population in some *relevant* respect, then we may start worrying about a lack of external validity.

That said, it is absolutely critical to realise that a study that uses only psychology students does not necessarily have a problem with external validity. I'll talk about this again later, but it's such a common mistake that I'm going to mention it here. The external validity of a study is threatened by the choice of population if (a) the population from which you sample your participants is very narrow (e.g., psychology students), and (b) the narrow population that you sampled from is systematically different from the general population *in some respect that is relevant to the psychological phenomenon that you intend to study*. The italicised part is the bit that lots of people forget. It is true that psychology undergraduates differ from the general population in lots of ways, and so a study that uses only psychology students *may* have problems with external validity. However, if those differences aren't very relevant to the phenomenon that you're studying, then there's nothing to worry about. To make this a bit more concrete here are two extreme examples:

- You want to measure "attitudes of the general public towards psychotherapy", but all of your participants are psychology students. This study would almost certainly have a problem with external validity.
- You want to measure the effectiveness of a visual illusion, and your participants are all psychology students. This study is unlikely to have a problem with external validity

Having just spent the last couple of paragraphs focusing on the choice of participants, since that's a big issue that everyone tends to worry most about, it's worth remembering that external validity is a broader concept. The following are also examples of things that might pose a threat to external validity, depending on what kind of study you're doing:

- People might answer a "psychology questionnaire" in a manner that doesn't reflect what they would do in real life.
- Your lab experiment on (say) "human learning" has a different structure to the learning problems people face in real life.

1.6.3 Construct validity

Construct validity is basically a question of whether you're measuring what you want to be mea-

suring. A measurement has good construct validity if it is actually measuring the correct theoretical construct, and bad construct validity if it doesn't. To give a very simple (if ridiculous) example, suppose I'm trying to investigate the rates with which university students cheat on their exams. And the way I attempt to measure it is by asking the cheating students to stand up in the lecture theatre so that I can count them. When I do this with a class of 300 students 0 people claim to be cheaters. So I therefore conclude that the proportion of cheaters in my class is 0%. Clearly this is a bit ridiculous. But the point here is not that this is a very deep methodological example, but rather to explain what construct validity is. The problem with my measure is that while I'm *trying* to measure "the proportion of people who cheat" what I'm actually measuring is "the proportion of people stupid enough to own up to cheating, or bloody minded enough to pretend that they do". Obviously, these aren't the same thing! So my study has gone wrong, because my measurement has very poor construct validity.

1.6.4 Face validity

Face validity simply refers to whether or not a measure "looks like" it's doing what it's supposed to, nothing more. If I design a test of intelligence, and people look at it and they say "no, that test doesn't measure intelligence", then the measure lacks face validity. It's as simple as that. Obviously, face validity isn't very important from a pure scientific perspective. After all, what we care about is whether or not the measure *actually* does what it's supposed to do, not whether it *looks like* it does what it's supposed to do. As a consequence, we generally don't care very much about face validity. That said, the concept of face validity serves three useful pragmatic purposes:

- Sometimes, an experienced scientist will have a "hunch" that a particular measure won't work. While these sorts of hunches have no strict evidentiary value, it's often worth paying attention to them. Because often times people have knowledge that they can't quite verbalise, so there might be something to worry about even if you can't quite say why. In other words, when someone you trust criticises the face validity of your study, it's worth taking the time to think more carefully about your design to see if you can think of reasons why it might go awry. Mind you, if you don't find any reason for concern, then you should probably not worry. After all, face validity really doesn't matter very much.

- Often (very often), completely uninformed people will also have a "hunch" that your research is crap. And they'll criticise it on the internet or something. On close inspection you may notice that these criticisms are actually focused entirely on how the study "looks", but not on anything deeper. The concept of face validity is useful for gently explaining to people that they need to substantiate their arguments further.

- Expanding on the last point, if the beliefs of untrained people are critical (e.g., this is often

the case for applied research where you actually want to convince policy makers of something or other) then you *have* to care about face validity. Simply because, whether you like it or not, a lot of people will use face validity as a proxy for real validity. If you want the government to change a law on scientific psychological grounds, then it won't matter how good your studies "really" are. If they lack face validity you'll find that politicians ignore you. Of course, it's somewhat unfair that policy often depends more on appearance than fact, but that's how things go.

1.6.5 Ecological validity

Ecological validity is a different notion of validity, which is similar to external validity, but less important. The idea is that, in order to be ecologically valid, the entire set up of the study should closely approximate the real world scenario that is being investigated. In a sense, ecological validity is a kind of face validity. It relates mostly to whether the study "looks" right, but with a bit more rigour to it. To be ecologically valid the study has to look right in a fairly specific way. The idea behind it is the intuition that a study that is ecologically valid is more likely to be externally valid. It's no guarantee, of course. But the nice thing about ecological validity is that it's much easier to check whether a study is ecologically valid than it is to check whether a study is externally valid. A simple example would be eyewitness identification studies. Most of these studies tend to be done in a university setting, often with a fairly simple array of faces to look at, rather than a line up. The length of time between seeing the "criminal" and being asked to identify the suspect in the "line up" is usually shorter. The "crime" isn't real so there's no chance of the witness being scared, and there are no police officers present so there's not as much chance of feeling pressured. These things all mean that the study *definitely* lacks ecological validity. They might (but might not) mean that it also lacks external validity.

1.7

Confounds, artefacts and other threats to validity

If we look at the issue of validity in the most general fashion the two biggest worries that we have are *confounders* and *artefacts*. These two terms are defined in the following way:

- **Confounder**: A confounder is an additional, often unmeasured variable[*6] that turns out to be related to both the predictors and the outcome. The existence of confounders threatens the internal validity of the study because you can't tell whether the predictor causes the outcome, or if the confounding variable causes it.
- **Artefact**: A result is said to be "artefactual" if it only holds in the special situation that you happened to test in your study. The possibility that your result is an artefact describes a threat to your external validity, because it raises the possibility that you can't generalise or apply your results to the actual population that you care about.

As a general rule confounders are a bigger concern for non-experimental studies, precisely because they're not proper experiments. By definition, you're leaving lots of things uncontrolled, so there's a lot of scope for confounders being present in your study. Experimental research tends to be much less vulnerable to confounders. The more control you have over what happens during the study, the more you can prevent confounders from affecting the results. With random allocation, for example, confounders are distributed randomly, and evenly, between different groups.

However, there are always swings and roundabouts and when we start thinking about artefacts rather than confounders the shoe is very firmly on the other foot. For the most part, artefactual results tend to be a concern for experimental studies than for non-experimental studies. To see this, it helps to realise that the reason that a lot of studies are non-experimental is precisely because what the researcher is trying to do is examine human behaviour in a more naturalistic context. By working in a more real-world context you lose experimental control (making yourself vulnerable to confounders), but because you tend to be studying human psychology "in the wild" you reduce the chances of getting an artefactual result. Or, to put it another way, when you take psychology out of the wild and bring it into the lab (which we usually have to do to gain our experimental control), you always run the risk of accidentally studying something different to what you wanted to study.

Be warned though. The above is a rough guide only. It's absolutely possible to have confounders in an experiment, and to get artefactual results with non-experimental studies. This can happen for all sorts of reasons, not least of which is experimenter or researcher error. In practice, it's really hard to think everything through ahead of time and even very good researchers make mistakes.

Although there's a sense in which almost any threat to validity can be characterised as a confounder or an artefact, they're pretty vague concepts. So let's have a look at some of the most common

---

[*6]The reason why I say that it's unmeasured is that if you *have* measured it, then you can use some fancy statistical tricks to deal with the confounder. Because of the existence of these statistical solutions to the problem of confounders, we often refer to a confounder that we have measured and dealt with as a *covariate*. Dealing with covariates is a more advanced topic, but I thought I'd mention it in passing since it's kind of comforting to at least know that this stuff exists.

examples.

1.7.1 History effects

History effects refer to the possibility that specific events may occur during the study that might influence the outcome measure. For instance, something might happen in between a pre-test and a post-test. Or in-between testing participant 23 and participant 24. Alternatively, it might be that you're looking at a paper from an older study that was perfectly valid for its time, but the world has changed enough since then that the conclusions are no longer trustworthy. Examples of things that would count as history effects are:

- You're interested in how people think about risk and uncertainty. You started your data collection in December 2010. But finding participants and collecting data takes time, so you're still finding new people in February 2011. Unfortunately for you (and even more unfortunately for others), the Queensland floods occurred in January 2011 causing billions of dollars of damage and killing many people. Not surprisingly, the people tested in February 2011 express quite different beliefs about handling risk than the people tested in December 2010. Which (if any) of these reflects the "true" beliefs of participants? I think the answer is probably both. The Queensland floods genuinely changed the beliefs of the Australian public, though possibly only temporarily. The key thing here is that the "history" of the people tested in February is quite different to people tested in December.

- You're testing the psychological effects of a new anti-anxiety drug. So what you do is measure anxiety before administering the drug (e.g., by self-report, and taking physiological measures). Then you administer the drug, and afterwards you take the same measures. In the middle however, because your lab is in Los Angeles, there's an earthquake which increases the anxiety of the participants.

1.7.2 Maturation effects

As with history effects, **maturational effects** are fundamentally about change over time. However, maturation effects aren't in response to specific events. Rather, they relate to how people change on their own over time. We get older, we get tired, we get bored, etc. Some examples of maturation effects are:

- When doing developmental psychology research you need to be aware that children grow up quite rapidly. So, suppose that you want to find out whether some educational trick helps

with vocabulary size among 3 year olds. One thing that you need to be aware of is that the vocabulary size of children that age is growing at an incredible rate (multiple words per day) all on its own. If you design your study without taking this maturational effect into account, then you won't be able to tell if your educational trick works.

- When running a very long experiment in the lab (say, something that goes for 3 hours) it's very likely that people will begin to get bored and tired, and that this maturational effect will cause performance to decline regardless of anything else going on in the experiment

1.7.3 Repeated testing effects

An important type of history effect is the effect of **repeated testing**. Suppose I want to take two measurements of some psychological construct (e.g., anxiety). One thing I might be worried about is if the first measurement has an effect on the second measurement. In other words, this is a history effect in which the "event" that influences the second measurement is the first measurement itself! This is not at all uncommon. Examples of this include:

- *Learning and practice*: e.g., "intelligence" at time 2 might appear to go up relative to time 1 because participants learned the general rules of how to solve "intelligence-test-style" questions during the first testing session.
- *Familiarity with the testing situation*: e.g., if people are nervous at time 1, this might make performance go down. But after sitting through the first testing situation they might calm down a lot precisely because they've seen what the testing looks like.
- *Auxiliary changes caused by testing*: e.g., if a questionnaire assessing mood is boring then mood rating at measurement time 2 is more likely to be "bored" precisely because of the boring measurement made at time 1.

1.7.4 Selection bias

Selection bias is a pretty broad term. Suppose that you're running an experiment with two groups of participants where each group gets a different "treatment", and you want to see if the different treatments lead to different outcomes. However, suppose that, despite your best efforts, you've ended up with a gender imbalance across groups (say, group A has 80% females and group B has 50% females). It might sound like this could never happen but, trust me, it can. This is an example of a selection bias, in which the people "selected into" the two groups have different characteristics. If any of those characteristics turns out to be relevant (say, your treatment works

better on females than males) then you're in a lot of trouble.

1.7.5 Differential attrition

When thinking about the effects of attrition, it is sometimes helpful to distinguish between two different types. The first is **homogeneous attrition**, in which the attrition effect is the same for all groups, treatments or conditions. In the example I gave above, the attrition would be homogeneous if (and only if) the easily bored participants are dropping out of all of the conditions in my experiment at about the same rate. In general, the main effect of homogeneous attrition is likely to be that it makes your sample unrepresentative. As such, the biggest worry that you'll have is that the generalisability of the results decreases. In other words, you lose external validity.

The second type of attrition is **heterogeneous attrition**, in which the attrition effect is different for different groups. More often called **differential attrition**, this is a kind of selection bias that is caused by the study itself. Suppose that, for the first time ever in the history of psychology, I manage to find the perfectly balanced and representative sample of people. I start running "Dani's incredibly long and tedious experiment" on my perfect sample but then, because my study is incredibly long and tedious, lots of people start dropping out. I can't stop this. Participants absolutely have the right to stop doing any experiment, any time, for whatever reason they feel like, and as researchers we are morally (and professionally) obliged to remind people that they do have this right. So, suppose that "Dani's incredibly long and tedious experiment" has a very high drop out rate. What do you suppose the odds are that this drop out is random? Answer: zero. Almost certainly the people who remain are more conscientious, more tolerant of boredom, etc., than those that leave. To the extent that (say) conscientiousness is relevant to the psychological phenomenon that I care about, this attrition can decrease the validity of my results.

Here's another example. Suppose I design my experiment with two conditions. In the "treatment" condition, the experimenter insults the participant and then gives them a questionnaire designed to measure obedience. In the "control" condition, the experimenter engages in a bit of pointless chitchat and then gives them the questionnaire. Leaving aside the questionable scientific merits and dubious ethics of such a study, let's have a think about what might go wrong here. As a general rule, when someone insults me to my face I tend to get much less co-operative. So, there's a pretty good chance that a lot more people are going to drop out of the treatment condition than the control condition. And this drop out isn't going to be random. The people most likely to drop out would probably be the people who don't care all that much about the importance of obediently sitting through the experiment. Since the most bloody minded and disobedient people all left the treatment group but not the control group, we've introduced a confound: the people who actually took the questionnaire in the treatment group were *already* more likely to be dutiful and obedient

than the people in the control group. In short, in this study insulting people doesn't make them more obedient. It makes the more disobedient people leave the experiment! The internal validity of this experiment is completely shot.

1.7.6 Non-response bias

Non-response bias is closely related to selection bias and to differential attrition. The simplest version of the problem goes like this. You mail out a survey to 1000 people but only 300 of them reply. The 300 people who replied are almost certainly not a random subsample. People who respond to surveys are systematically different to people who don't. This introduces a problem when trying to generalise from those 300 people who replied to the population at large, since you now have a very non-random sample. The issue of non-response bias is more general than this, though. Among the (say) 300 people that did respond to the survey, you might find that not everyone answers every question. If (say) 80 people chose not to answer one of your questions, does this introduce problems? As always, the answer is maybe. If the question that wasn't answered was on the last page of the questionnaire, and those 80 surveys were returned with the last page missing, there's a good chance that the missing data isn't a big deal; probably the pages just fell off. However, if the question that 80 people didn't answer was the most confrontational or invasive personal question in the questionnaire, then almost certainly you've got a problem. In essence, what you're dealing with here is what's called the problem of **missing data**. If the data that is missing was "lost" randomly, then it's not a big problem. If it's missing systematically, then it can be a big problem.

1.7.7 Regression to the mean

Regression to the mean refers to any situation where you select data based on an extreme value on some measure. Because the variable has natural variation it almost certainly means that when you take a subsequent measurement the later measurement will be less extreme than the first one, purely by chance.

Here's an example. Suppose I'm interested in whether a psychology education has an adverse effect on very smart kids. To do this, I find the 20 psychology I students with the best high school grades and look at how well they're doing at university. It turns out that they're doing a lot better than average, but they're not topping the class at university even though they did top their classes at high school. What's going on? The natural first thought is that this must mean that the psychology classes must be having an adverse effect on those students. However, while that might very well be the explanation, it's more likely that what you're seeing is an example of "regression to the mean". To see how it works, let's take a moment to think about what is required to get the best mark in

a class, regardless of whether that class be at high school or at university. When you've got a big class there are going to be *lots* of very smart people enrolled. To get the best mark you have to be very smart, work very hard, and be a bit lucky. The exam has to ask just the right questions for your idiosyncratic skills, and you have to avoid making any dumb mistakes (we all do that sometimes) when answering them. And that's the thing, whilst intelligence and hard work are transferable from one class to the next, luck isn't. The people who got lucky in high school won't be the same as the people who get lucky at university. That's the very definition of "luck". The consequence of this is that when you select people at the very extreme values of one measurement (the top 20 students), you're selecting for hard work, skill and luck. But because the luck doesn't transfer to the second measurement (only the skill and work), these people will all be expected to drop a little bit when you measure them a second time (at university). So their scores fall back a little bit, back towards everyone else. This is regression to the mean.

Regression to the mean is surprisingly common. For instance, if two very tall people have kids their children will tend to be taller than average but not as tall as the parents. The reverse happens with very short parents. Two very short parents will tend to have short children, but nevertheless those kids will tend to be taller than the parents. It can also be extremely subtle. For instance, there have been studies done that suggested that people learn better from negative feedback than from positive feedback. However, the way that people tried to show this was to give people positive reinforcement whenever they did good, and negative reinforcement when they did bad. And what you see is that after the positive reinforcement people tended to do worse, but after the negative reinforcement they tended to do better. But notice that there's a selection bias here! When people do very well, you're selecting for "high" values, and so you should *expect*, because of regression to the mean, that performance on the next trial should be worse regardless of whether reinforcement is given. Similarly, after a bad trial, people will tend to improve all on their own. The apparent superiority of negative feedback is an artefact caused by regression to the mean (**Kahneman1973**).

1.7.8 Experimenter bias

Experimenter bias can come in multiple forms. The basic idea is that the experimenter, despite the best of intentions, can accidentally end up influencing the results of the experiment by subtly communicating the "right answer" or the "desired behaviour" to the participants. Typically, this occurs because the experimenter has special knowledge that the participant does not, for example the right answer to the questions being asked or knowledge of the expected pattern of performance for the condition that the participant is in. The classic example of this happening is the case study of "Clever Hans", which dates back to 1907 (**Pfungst1911**; **Hothersall2004**). Clever Hans was a horse that apparently was able to read and count and perform other human like feats of intelligence.

After Clever Hans became famous, psychologists started examining his behaviour more closely. It turned out that, not surprisingly, Hans didn't know how to do maths. Rather, Hans was responding to the human observers around him, because the humans did know how to count and the horse had learned to change its behaviour when people changed theirs.

The general solution to the problem of experimenter bias is to engage in double blind studies, where neither the experimenter nor the participant knows which condition the participant is in or knows what the desired behaviour is. This provides a very good solution to the problem, but it's important to recognise that it's not quite ideal, and hard to pull off perfectly. For instance, the obvious way that I could try to construct a double blind study is to have one of my Ph.D. students (one who doesn't know anything about the experiment) run the study. That feels like it should be enough. The only person (me) who knows all the details (e.g., correct answers to the questions, assignments of participants to conditions) has no interaction with the participants, and the person who does all the talking to people (the Ph.D. student) doesn't know anything. Except for the reality that the last part is very unlikely to be true. In order for the Ph.D. student to run the study effectively they need to have been briefed by me, the researcher. And, as it happens, the Ph.D. student also knows me and knows a bit about my general beliefs about people and psychology (e.g., I tend to think humans are much smarter than psychologists give them credit for). As a result of all this, it's almost impossible for the experimenter to avoid knowing a little bit about what expectations I have. And even a little bit of knowledge can have an effect. Suppose the experimenter accidentally conveys the fact that the participants are expected to do well in this task. Well, there's a thing called the "Pygmalion effect", where if you expect great things of people they'll tend to rise to the occasion. But if you expect them to fail then they'll do that too. In other words, the expectations become a self-fulfilling prophesy.

1.7.9 Demand effects and reactivity

When talking about experimenter bias, the worry is that the experimenter's knowledge or desires for the experiment are communicated to the participants, and that these can change people's behaviour (**Rosenthal1966**). However, even if you manage to stop this from happening, it's almost impossible to stop people from knowing that they're part of a psychological study. And the mere fact of knowing that someone is watching or studying you can have a pretty big effect on behaviour. This is generally referred to as **reactivity** or **demand effects**. The basic idea is captured by the Hawthorne effect: people alter their performance because of the attention that the study focuses on them. The effect takes its name from a study that took place in the "Hawthorne Works" factory outside of Chicago (**Adair1984**). This study, from the 1920s, looked at the effects of factory lighting on worker productivity. But, importantly, change in worker behaviour occurred because the workers *knew* they

were being studied, rather than any effect of factory lighting.

To get a bit more specific about some of the ways in which the mere fact of being in a study can change how people behave, it helps to think like a social psychologist and look at some of the *roles* that people might *adopt* during an experiment but might *not adopt* if the corresponding events were occurring in the real world:

- The *good participant* tries to be too helpful to the researcher. He or she seeks to figure out the experimenter's hypotheses and confirm them.
- The *negative participant* does the exact opposite of the good participant. He or she seeks to break or destroy the study or the hypothesis in some way.
- The *faithful participant* is unnaturally obedient. He or she seeks to follow instructions perfectly, regardless of what might have happened in a more realistic setting.
- The *apprehensive participant* gets nervous about being tested or studied, so much so that his or her behaviour becomes highly unnatural, or overly socially desirable.

1.7.10 Placebo effects

The **placebo effect** is a specific type of demand effect that we worry a lot about. It refers to the situation where the mere fact of being treated causes an improvement in outcomes. The classic example comes from clinical trials. If you give people a completely chemically inert drug and tell them that it's a cure for a disease, they will tend to get better faster than people who aren't treated at all. In other words, it is people's belief that they are being treated that causes the improved outcomes, not the drug.

However, the current consensus in medicine is that true placebo effects are quite rare and most of what was previously considered placebo effect is in fact some combination of natural healing (some people just get better on their own), regression to the mean and other quirks of study design. Of interest to psychology is that the strongest evidence for at least some placebo effect is in self-reported outcomes, most notably in treatment of pain (**hrobjartsson2010**).

1.7.11 Situation, measurement and sub-population effects

In some respects, these terms are a catch-all term for "all other threats to external validity". They refer to the fact that the choice of sub-population from which you draw your participants, the location, timing and manner in which you run your study (including who collects the data) and the tools that you use to make your measurements might all be influencing the results. Specifically, the worry is that these things might be influencing the results in such a way that the results won't

generalise to a wider array of people, places and measures.

1.7.12 Fraud, deception and self-deception

> *It is difficult to get a man to understand something, when his salary depends on his not understanding it.*
>   – Upton Sinclair

There's one final thing I feel I should mention. While reading what the textbooks often have to say about assessing the validity of a study I couldn't help but notice that they seem to make the assumption that the researcher is honest. I find this hilarious. While the vast majority of scientists are honest, in my experience at least, some are not.[*7] Not only that, as I mentioned earlier, scientists are not immune to belief bias. It's easy for a researcher to end up deceiving themselves into believing the wrong thing, and this can lead them to conduct subtly flawed research and then hide those flaws when they write it up. So you need to consider not only the (probably unlikely) possibility of outright fraud, but also the (probably quite common) possibility that the research is unintentionally "slanted". I opened a few standard textbooks and didn't find much of a discussion of this problem, so here's my own attempt to list a few ways in which these issues can arise:

- **Data fabrication**. Sometimes, people just make up the data. This is occasionally done with "good" intentions. For instance, the researcher believes that the fabricated data do reflect the truth, and may actually reflect "slightly cleaned up" versions of actual data. On other occasions, the fraud is deliberate and malicious. Some high-profile examples where data fabrication has been alleged or shown include Cyril Burt (a psychologist who is thought to have fabricated some of his data), Andrew Wakefield (who has been accused of fabricating his data connecting the MMR vaccine to autism) and Hwang Woo-suk (who falsified a lot of his data on stem cell research).
- **Hoaxes**. Hoaxes share a lot of similarities with data fabrication, but they differ in the intended purpose. A hoax is often a joke, and many of them are intended to be (eventually) discovered. Often, the point of a hoax is to discredit someone or some field. There's quite a few well known scientific hoaxes that have occurred over the years (e.g., Piltdown man) and some were deliberate attempts to discredit particular fields of research (e.g., the Sokal affair).
- **Data misrepresentation**. While fraud gets most of the headlines, it's much more common

---

[*7]Some people might argue that if you're not honest then you're not a real scientist. Which does have some truth to it I guess, but that's disingenuous (look up the "No true Scotsman" fallacy). The fact is that there are lots of people who are employed ostensibly as scientists, and whose work has all of the trappings of science, but who are outright fraudulent. Pretending that they don't exist by saying that they're not scientists is just muddled thinking.

in my experience to see data being misrepresented. When I say this I'm not referring to newspapers getting it wrong (which they do, almost always). I'm referring to the fact that often the data don't actually say what the researchers think they say. My guess is that, almost always, this isn't the result of deliberate dishonesty but instead is due to a lack of sophistication in the data analyses. For instance, think back to the example of Simpson's paradox that I discussed in the beginning of this book. It's very common to see people present "aggregated" data of some kind and sometimes, when you dig deeper and find the raw data yourself you find that the aggregated data tell a different story to the disaggregated data. Alternatively, you might find that some aspect of the data is being hidden, because it tells an inconvenient story (e.g., the researcher might choose not to refer to a particular variable). There's a lot of variants on this, many of which are very hard to detect.

- **Study "misdesign"**. Okay, this one is subtle. Basically, the issue here is that a researcher designs a study that has built-in flaws and those flaws are never reported in the paper. The data that are reported are completely real and are correctly analysed, but they are produced by a study that is actually quite wrongly put together. The researcher really wants to find a particular effect and so the study is set up in such a way as to make it "easy" to (artefactually) observe that effect. One sneaky way to do this, in case you're feeling like dabbling in a bit of fraud yourself, is to design an experiment in which it's obvious to the participants what they're "supposed" to be doing, and then let reactivity work its magic for you. If you want you can add all the trappings of double blind experimentation but it won't make a difference since the study materials themselves are subtly telling people what you want them to do. When you write up the results the fraud won't be obvious to the reader. What's obvious to the participant when they're in the experimental context isn't always obvious to the person reading the paper. Of course, the way I've described this makes it sound like it's always fraud. Probably there are cases where this is done deliberately, but in my experience the bigger concern has been with unintentional misdesign. The researcher *believes* and so the study just happens to end up with a built in flaw, and that flaw then magically erases itself when the study is written up for publication.

- **Data mining & post hoc hypothesising**. Another way in which the authors of a study can more or less misrepresent the data is by engaging in what's referred to as "data mining" (**Gelman2014**). As we'll discuss later, if you keep trying to analyse your data in lots of different ways, you'll eventually find something that "looks" like a real effect but isn't. This is referred to as "data mining". It used to be quite rare because data analysis used to take weeks, but now that everyone has very powerful statistical software on their computers it's becoming very common. Data mining per se isn't "wrong", but the more that you do it the bigger the risk you're taking. The thing that is wrong, and I suspect is very common, is

unacknowledged data mining. That is, the researcher runs every possible analysis known to humanity, finds the one that works, and then pretends that this was the only analysis that they ever conducted. Worse yet, they often "invent" a hypothesis after looking at the data to cover up the data mining. To be clear. It's not wrong to change your beliefs after looking at the data, and to reanalyse your data using your new "post hoc" hypotheses. What is wrong (and I suspect common) is failing to acknowledge that you've done. If you acknowledge that you did it then other researchers are able to take your behaviour into account. If you don't, then they can't. And that makes your behaviour deceptive. Bad!

- **Publication bias & self-censoring**. Finally, a pervasive bias is "non-reporting" of negative results. This is almost impossible to prevent. Journals don't publish every article that is submitted to them. They prefer to publish articles that find "something". So, if 20 people run an experiment looking at whether reading *Finnegans Wake* causes insanity in humans, and 19 of them find that it doesn't, which one do you think is going to get published? Obviously, it's the one study that did find that *Finnegans Wake* causes insanity.[*8] This is an example of a *publication bias*. Since no-one ever published the 19 studies that didn't find an effect, a naive reader would never know that they existed. Worse yet, most researchers "internalise" this bias and end up *self-censoring* their research. Knowing that negative results aren't going to be accepted for publication, they never even try to report them. As a friend of mine says "for every experiment that you get published, you also have 10 failures". And she's right. The catch is, while some (maybe most) of those studies are failures for boring reasons (e.g. you stuffed something up) others might be genuine "null" results that you ought to acknowledge when you write up the "good" experiment. And telling which is which is often hard to do. A good place to start is a paper by **Ioannidis2005** with the depressing title "Why most published research findings are false". I'd also suggest taking a look at work by **Kuhberger2014** presenting statistical evidence that this actually happens in psychology.

There's probably a lot more issues like this to think about, but that'll do to start with. What I really want to point out is the blindingly obvious truth that real world science is conducted by actual humans, and only the most gullible of people automatically assumes that everyone else is honest and impartial. Actual scientists aren't usually *that* naive, but for some reason the world likes to pretend that we are, and the textbooks we usually write seem to reinforce that stereotype.

1.8 _____

[*8] Clearly, the real effect is that only insane people would even try to read *Finnegans Wake*.

Summary

This chapter isn't really meant to provide a comprehensive discussion of psychological research methods. It would require another volume just as long as this one to do justice to the topic. However, in real life statistics and study design are so tightly intertwined that it's very handy to discuss some of the key topics. In this chapter, I've briefly discussed the following topics:

- *Introduction to psychological measurement* (Section **??**). What does it mean to operationalise a theoretical construct? What does it mean to have variables and take measurements?
- *Scales of measurement and types of variables* (Section **??**). Remember that there are *two* different distinctions here. There's the difference between discrete and continuous data, and there's the difference between the four different scale types (nominal, ordinal, interval and ratio).
- *Reliability of a measurement* (Section **??**). If I measure the "same" thing twice, should I expect to see the same result? Only if my measure is reliable. But what does it mean to talk about doing the "same" thing? Well, that's why we have different types of reliability. Make sure you remember what they are.
- *Terminology: predictors and outcomes* (Section **??**). What roles do variables play in an analysis? Can you remember the difference between predictors and outcomes? Dependent and independent variables? Etc.
- *Experimental and non-experimental research designs* (Section **??**). What makes an experiment an experiment? Is it a nice white lab coat, or does it have something to do with researcher control over variables?
- *Validity and its threats* (Section **??**). Does your study measure what you want it to? How might things go wrong? And is it my imagination, or was that a very long list of possible ways in which things can go wrong?

All this should make clear to you that study design is a critical part of research methodology. I built this chapter from the classic little book by **Campbell1963**, but there are of course a large number of textbooks out there on research design. Spend a few minutes with your favourite search engine and you'll find dozens.